

Problem statement

Data collection

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\4_drug200 - 4_drug200.csv")
data
```

```
Out[2]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY
...
195	56	F	LOW	HIGH	11.567	drugC
196	16	M	LOW	HIGH	12.006	drugC
197	52	M	NORMAL	HIGH	9.894	drugX
198	23	M	NORMAL	NORMAL	14.020	drugX
199	40	F	LOW	NORMAL	11.349	drugX

200 rows × 6 columns

head

```
In [3]: # to display first 8 dataset values
da=data.head(8)
da
```

Out[3]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY
5	22	F	NORMAL	HIGH	8.607	drugX
6	49	F	NORMAL	HIGH	16.275	drugY
7	41	M	LOW	HIGH	11.037	drugC

info

In [4]:

to identify missing values
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
Column Non-Null Count Dtype
--- -
0 Age 200 non-null int64
1 Sex 200 non-null object
2 BP 200 non-null object
3 Cholesterol 200 non-null object
4 Na_to_K 200 non-null float64
5 Drug 200 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB

describe

In [5]:

to display summary of the dataset
data.describe()

Out[5]:

	Age	Na_to_K
count	200.000000	200.000000
mean	44.315000	16.084485
std	16.544315	7.223956
min	15.000000	6.269000
25%	31.000000	10.445500
50%	45.000000	13.936500
75%	58.000000	19.380000

	Age	Na_to_K
max	74.000000	38.247000

columns

```
In [6]: # to display headings of the dataset
data.columns
```

```
Out[6]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
In [7]: a=data.dropna(axis=1)
a
```

```
Out[7]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY
...
195	56	F	LOW	HIGH	11.567	drugC
196	16	M	LOW	HIGH	12.006	drugC
197	52	M	NORMAL	HIGH	9.894	drugX
198	23	M	NORMAL	NORMAL	14.020	drugX
199	40	F	LOW	NORMAL	11.349	drugX

200 rows × 6 columns


```
In [8]: a.columns
```

```
Out[8]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

EDA and Visualization

```
In [9]: sns.pairplot(a)
```


```
Out[9]: <seaborn.axisgrid.PairGrid at 0x2200e901e20>
```



distribution plot

```
In [10]: sns.displot(a["Age"])
```

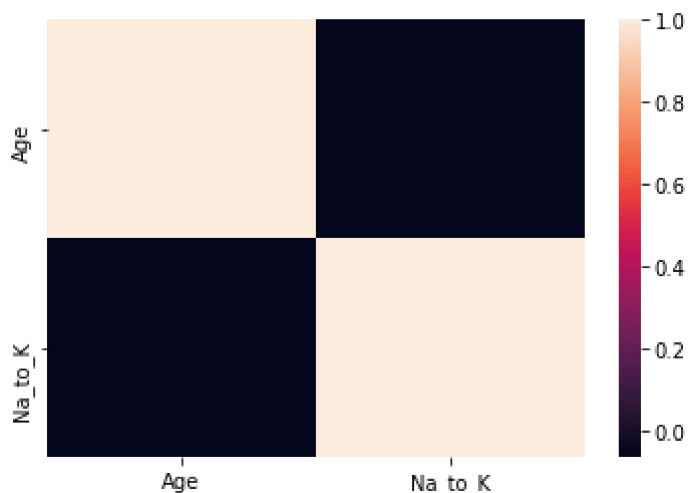
```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x22010243e20>
```



correlation

```
In [11]: dat=data[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug']]
          sns.heatmap(dat.corr())
```

Out[11]: <AxesSubplot:>



To train the model-Model Building

```
In [12]: x=a[['Age']]
         y=a['Age']
```

```
In [13]: # to split my dataset into training and test data
         from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
         lr= LinearRegression()
         lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: print(lr.intercept_)
```

-7.105427357601002e-15


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
         coeff
```

```
Out[16]:
```

	Co-efficient
Age	1.0

```
In [17]: prediction=lr.predict(x_test)
         plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x22010b8dc10>



```
In [18]: print(lr.score(x_test,y_test))
```

1.0

```
In [19]: lr.score(x_train,y_train)
```

Out[19]: 1.0

Ridge regression

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

Out[21]: 0.9999999252234555

```
In [22]: rr.score(x_train,y_train)
```

Out[22]: 0.9999999271077268

Lasso regression

```
In [23]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
la.score(x_train,y_train)
```

Out[23]: 0.9985705396815336

```
In [24]: la.score(x_test,y_test)
```

Out[24]: 0.9985335880131809

In [25]: `lr.score(x_train,y_train)`

Out[25]: 1.0

Ridge regression

In [26]: `from sklearn.linear_model import Ridge,Lasso`

In [27]: `rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)`

Out[27]: 0.9999999252234555

In [28]: `rr.score(x_train,y_train)`

Out[28]: 0.9999999271077268

Lasso regression

In [29]: `la=Lasso(alpha=10)
la.fit(x_train,y_train)
la.score(x_train,y_train)`

Out[29]: 0.9985705396815336

In [30]: `la.score(x_test,y_test)`

Out[30]: 0.9985335880131809

In [31]: `from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)`

Out[31]: ElasticNet()

In [32]: `print(en.coef_)`

[0.99622631]

In [33]: `print(en.intercept_)`

0.16415536806922404

```
In [34]: predict=en.predict(x_test)
```

```
In [35]: print(en.score(x_test,y_test))
```

0.9999853911657179

```
In [36]: from sklearn import metrics
```

```
In [37]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,predict))
```

Mean Absolute error: 0.056919773219404404

```
In [38]: print("Mean Squared error:",metrics.mean_squared_error(y_test,predict))
```

Mean Squared error: 0.004170866825641406

```
In [39]: print("Root squared error:",np.sqrt(metrics.mean_squared_error(y_test,predict)))
```

Root squared error: 0.0645822485334895