

Problem statement

Data collection

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\PLACEMENT.csv")
data
```

```
Out[2]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26	1
1	7.46	38	1
2	7.54	40	1
3	6.42	8	1
4	7.23	17	0
...
995	8.87	44	1
996	9.12	65	1
997	4.89	34	0
998	8.62	46	1
999	4.90	10	1

1000 rows × 3 columns

head

```
In [3]: # to display first 8 dataset values
da=data.head(8)
da
```

Out[3]:

	cgpa	placement_exam_marks	placed
0	7.19	26	1
1	7.46	38	1
2	7.54	40	1
3	6.42	8	1
4	7.23	17	0
5	7.30	23	1
6	6.69	11	0
7	7.12	39	1

info

```
In [4]: # to identify missing values
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cgpa                   1000 non-null   float64
1   placement_exam_marks  1000 non-null   int64
2   placed                 1000 non-null   int64
dtypes: float64(1), int64(2)
memory usage: 23.6 KB
```

describe

```
In [5]: # to display summary of the dataset
data.describe()

Out[5]:
```

	cgpa	placement_exam_marks	placed
count	1000.000000	1000.000000	1000.000000
mean	6.961240	32.225000	0.489000
std	0.615898	19.130822	0.500129
min	4.890000	0.000000	0.000000
25%	6.550000	17.000000	0.000000
50%	6.960000	28.000000	0.000000
75%	7.370000	44.000000	1.000000
max	9.120000	100.000000	1.000000

columns

```
In [6]: # to display headings of the dataset  
data.columns
```

```
Out[6]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

```
In [7]: a=data.dropna(axis=1)  
a  
b=a.head(8)  
b
```

```
Out[7]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26	1
1	7.46	38	1
2	7.54	40	1
3	6.42	8	1
4	7.23	17	0
5	7.30	23	1
6	6.69	11	0
7	7.12	39	1

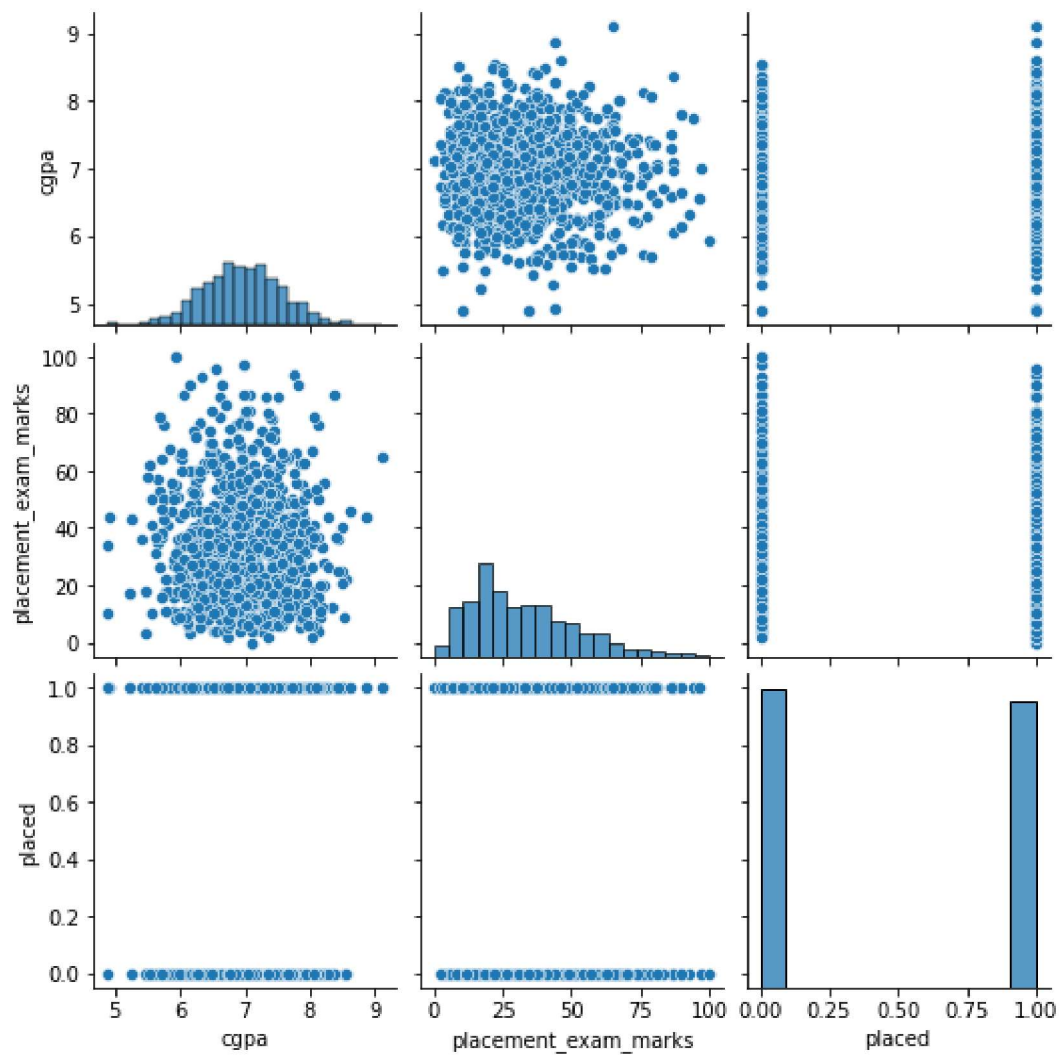
```
In [8]: a.columns
```

```
Out[8]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

EDA and Visualization

```
In [9]: sns.pairplot(data)
```

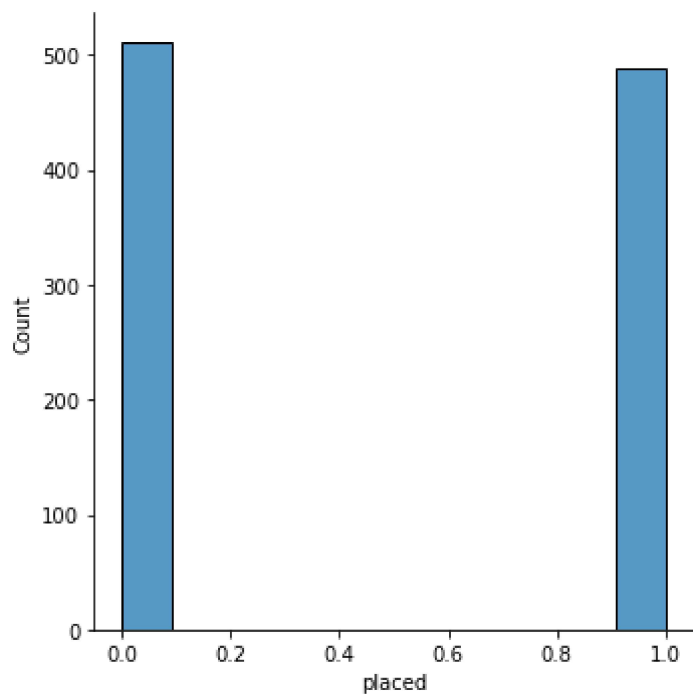
```
Out[9]: <seaborn.axisgrid.PairGrid at 0x13c55f41550>
```



distribution plot

```
In [10]: sns.displot(a["placed"])
```

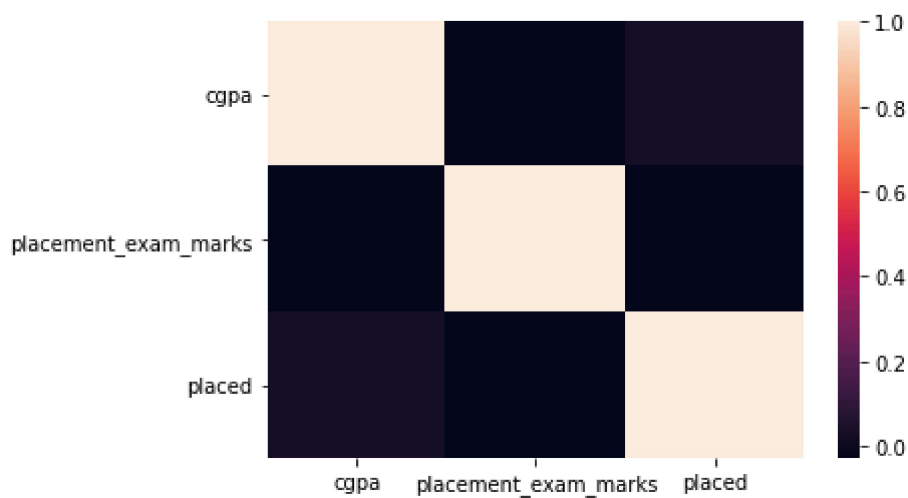
```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x13c57bcd20>
```



correlation

```
In [11]: dat=data[['cgpa', 'placement_exam_marks', 'placed']]
          sns.heatmap(dat.corr())
```

Out[11]: <AxesSubplot:>



To train the model-Model Building

```
In [12]: x=a[['placed']]
          y=a['placement_exam_marks']
```

```
In [13]: # to split my dataset into training and test data
          from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
lr= LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: print(lr.intercept_)

33.328611898017
```


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

```
Out[16]:
```

	Co-efficient
placed	-1.602387

```
In [17]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x13c58676370>



```
In [18]: print(lr.score(x_test,y_test))

-0.0059288704050051155
```

```
In [19]: lr.score(x_train,y_train)
```

Out[19]: 0.0017389154837557097

Ridge regression

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
         rr.score(x_test,y_test)
```

```
Out[21]: -0.005677495470711946
```

```
In [22]: rr.score(x_train,y_train)
```

```
Out[22]: 0.0017338339433277117
```

Lasso regression

```
In [23]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
         la.score(x_train,y_train)
```

```
Out[23]: 0.0
```

```
In [24]: la.score(x_test,y_test)
```

```
Out[24]: -0.0029795159037266927
```

```
In [25]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[25]: ElasticNet()
```

```
In [26]: print(en.coef_)
```

```
[-0.]
```

```
In [27]: print(en.intercept_)
```

```
32.534285714285716
```

```
In [34]: predict=en.predict(x_test)
```

```
In [31]: print(en.score(x_test,y_test))
```

```
-0.0029795159037266927
```

```
In [32]: from sklearn import metrics
```

```
In [35]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,predict))
```

Mean Absolute error: 15.382609523809524

In [36]:

```
print("Mean Squared error:", metrics.mean_squared_error(y_test, predict))
```

Mean Squared error: 357.78618503401356

In [38]:

```
print("Root squared error:", np.sqrt(metrics.mean_squared_error(y_test, predict)))
```

Root squared error: 18.915236848477832