

# Problem statement

## Data collection

## Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

## Importing dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\wine.csv")
data
```

```
Out[2]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...	...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	

1599 rows × 12 columns



## head

```
In [3]: # to display first 8 dataset values
data.head(8)
```

da

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qualit
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	

## info

In [4]:

```
# to identify missing values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   fixed acidity         1599 non-null   float64
 1   volatile acidity      1599 non-null   float64
 2   citric acid           1599 non-null   float64
 3   residual sugar        1599 non-null   float64
 4   chlorides             1599 non-null   float64
 5   free sulfur dioxide    1599 non-null   float64
 6   total sulfur dioxide   1599 non-null   float64
 7   density               1599 non-null   float64
 8   pH                   1599 non-null   float64
 9   sulphates             1599 non-null   float64
10   alcohol               1599 non-null   float64
11   quality               1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

## describe

In [5]:

```
# to display summary of the dataset
data.describe()
```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	quality
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.71
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.17
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.49
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.64
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.68
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.76
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.00

## columns


```
In [6]: # to display headings of the dataset
data.columns
```

```
Out[6]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
              'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
              'pH', 'sulphates', 'alcohol', 'quality'],
              dtype='object')
```

```
In [7]: a=data.dropna(axis=1)
a
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...	...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	

1599 rows × 12 columns



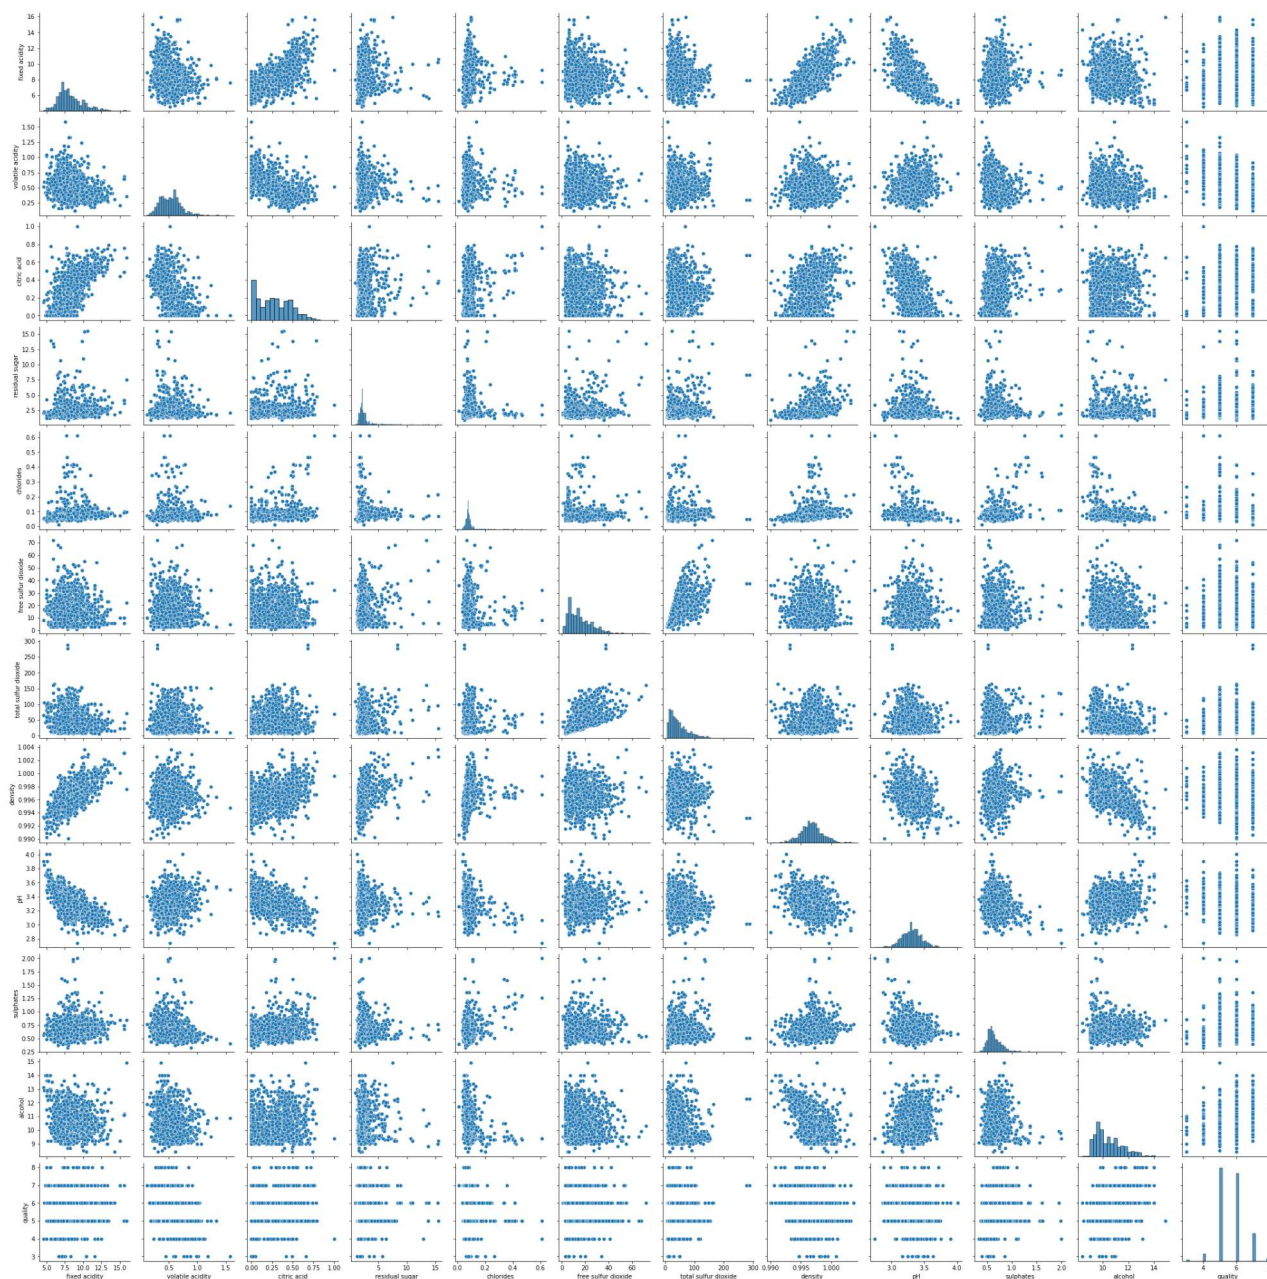
```
In [8]: a.columns
```

```
Out[8]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
             'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
             'pH', 'sulphates', 'alcohol', 'quality'],  
            dtype='object')
```

## EDA and Visualization

```
In [9]: sns.pairplot(a)
```


```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1cb97eafbb0>
```



distribution plot

```
In [10]: sns.displot(a["alcohol"])
```


```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x1cb9cd7b730>
```



## correlation

```
In [11]: dat=data[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
                  'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
                  'pH', 'sulphates', 'alcohol', 'quality']]
sns.heatmap(dat.corr())
```

```
Out[11]: <AxesSubplot:>
```



# To train the model-Model Building

```
In [12]: x=a[['quality']]
         y=a['quality']
```

```
In [13]: # to split my dataset into training and test data
         from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
         lr= LinearRegression()
         lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: print(lr.intercept_)
```

-3.552713678800501e-15


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
         coeff
```

Out[16]:

	Co-efficient
quality	1.0

```
In [17]: prediction=lr.predict(x_test)
         plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x1cba074c280>



```
In [18]: print(lr.score(x_test,y_test))
```



1.0

```
In [19]: lr.score(x_train,y_train)
```

Out[19]: 1.0

## Ridge regression

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)  
rr.score(x_test,y_test)
```

Out[21]: 0.9998276408763603

```
In [22]: rr.score(x_train,y_train)
```

Out[22]: 0.9998293261823417

## Lasso regression

```
In [23]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)  
la.score(x_train,y_train)
```

Out[23]: 0.0

```
In [24]: la.score(x_test,y_test)
```

Out[24]: -0.009874425992380198

```
In [25]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[25]: ElasticNet()

```
In [26]: print(en.coef_)
```

[0.14901675]

```
In [27]: print(en.intercept_)
```

4.77660926100045

```
In [28]: predict=en.predict(x_test)
```

```
In [29]: print(en.score(x_test,y_test))
```

0.2686767288051497

```
In [30]: from sklearn import metrics
```

```
In [31]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,predict))
```

Mean Absolute error: 0.5655334382540507

```
In [33]: print("Mean Squared error:",metrics.mean_squared_error(y_test,predict))
```

Mean Squared error: 0.4338391205887193

```
In [34]: print("Root squared error:",np.sqrt(metrics.mean_squared_error(y_test,predict)))
```

Root squared error: 0.6586646495666207

```
In [ ]:
```